

People Management Interview Questions & STAR Responses

Tailored for Director-Level Position | Yugandhar Reddy

1. Building and Scaling Teams

Q: Tell me about a time you built a team from scratch or significantly scaled an organization.

Situation: When I joined NetApp's Engineering Performance & Sustaining (EPS) group, there was no dedicated team to handle product quality KPIs, ONTAP patch management, or availability analytics. The organization relied on ad-hoc efforts spread across multiple groups with no unified ownership.

Task: I needed to bootstrap a team from zero that could own 24x7 SRE operations, define product quality KPIs, manage ONTAP monthly patches, and drive availability analytics — all while establishing credibility with senior leadership.

Action: I developed a hiring plan aligned to critical capability gaps, prioritizing SRE expertise, data analytics skills, and release engineering talent. I personally recruited and mentored engineers, establishing clear role definitions and career ladders. I created an operating model with defined KPIs, built an incident command control culture, and implemented a Correction of Errors (COE) process for blameless learning. As the team matured, I expanded it to 60+ engineers covering 24x7 operations across multiple time zones.

Result: The team grew from 0 to 60+ engineers. We reduced P1 disruption incidents by 60% through HA cluster automation, achieved 99.995% availability telemetry via AutoSupport infrastructure, and established a predictable half-yearly release model focused on Customer First principles. This team became the backbone of NetApp's product quality and reliability operations.

2. Managing Managers (Manager of Managers)

Q: Describe your experience managing other managers or directors. How do you develop leaders who report to you?

Situation: As Senior Director of HCI Engineering at NetApp, I managed a \$250M software product portfolio with 180+ SDEs distributed across USA, EMEA, and APAC — with 2 Directors reporting directly to me. Each Director owned distinct product lines and geographically dispersed teams with different maturity levels.

Task: I needed to ensure both Directors operated autonomously while maintaining alignment on quality standards, delivery cadence, and customer

outcomes. I also needed to develop their strategic thinking and P&L ownership capabilities.

Action: I established a leadership operating model where each Director owned their P&L, hiring plans, and delivery roadmaps. I held weekly 1:1s focused on strategic decisions rather than tactical execution, coached them on executive communication, and gave them increasing exposure to VP-level reviews. I implemented a shared "left-shift QA" paradigm across both organizations, creating a common engineering excellence standard. I also created cross-pollination opportunities — having their teams collaborate on shared infrastructure like the AWS-based SaaS log streaming service.

Result: Both Directors grew into fully autonomous leaders managing \$100M+ portfolios independently. We achieved 90% reduction in field regressions through the unified CI/CD culture, maintained 5-nines availability for cloud operations, and successfully served 8K+ HCI cluster deployments. One Director went on to lead an expanded organization, demonstrating the leadership pipeline I built.

3. Driving Organizational Change

Q: Tell me about a time you drove a significant cultural or process change across your organization.

Situation: When I took over as Lead SDM for Aurora MySQL at AWS, the operational burden on the team was unsustainable. Engineers were spending excessive time on manual operational tasks — toil that was consuming capacity that should have been directed toward feature development for a multi-billion dollar service.

Task: I needed to fundamentally shift the team's operating model from reactive operations to proactive automation, reducing operational burden while simultaneously delivering high-impact features like Serverless V2 and Zero Downtime upgrades.

Action: I championed an AI/ML infrastructure initiative (internally called "Minions") to automate repetitive operational tasks across global regions. I established a DevOps and Canary automation culture, implementing CI/CD infrastructure that shifted quality left in the development cycle. I restructured team priorities to dedicate capacity for automation investments, even when there was pressure to ship features. I also optimized our canary infrastructure to reduce unnecessary compute spend.

Result: The AI/ML Minions infrastructure reduced 56% of Aurora MySQL's operational burden, freeing engineers to focus on innovation. Canary optimizations saved \$6M per annum in CapEx. The team successfully delivered Aurora Serverless V2 (~\$1B revenue impact), FastRestart (24x improvement), and Parallel Export (100x speed improvement) — all while maintaining service reliability for millions of databases.

4. Hiring and Talent Strategy

Q: How do you approach hiring for senior technical roles, and can you give an example of a critical hire that transformed your team?

Situation: At AWS Aurora, I was leading 50+ SDEs across 3 SDMs, and we were launching Aurora Limitless (scale-out/multi-shard architecture) — a technically complex, greenfield initiative requiring deep distributed systems expertise that didn't exist on the current team.

Task: I needed to hire adept leaders and principal engineers who could architect and deliver a multi-shard scale-out database service while maintaining the operational excellence bar for the existing Aurora MySQL fleet.

Action: I defined the technical bar explicitly — candidates needed distributed systems depth, database internals knowledge, and the ability to operate in ambiguity. I personally participated in hiring loops for senior roles, calibrated interviewers on what "great" looked like for this problem space, and created a compelling narrative about the impact of Aurora Limitless to attract top talent. I also invested in internal talent development — identifying high-potential engineers on existing teams and creating stretch assignments on the Limitless project.

Result: We successfully staffed the Aurora Limitless initiative and moved it to Beta. The combination of external hires and internal talent development created a team capable of delivering a multi-shard architecture while the broader team continued delivering Zero Downtime upgrades for millions of databases in under 10 seconds. The hiring bar I set became the template for subsequent senior hiring across the Aurora organization.

5. Performance Management and Difficult Conversations

Q: Tell me about a time you had to manage underperformance or make a difficult people decision.

Situation: At NetApp, within my 180+ person HCI organization, I had a team that was consistently missing delivery milestones for a critical cloud integration (Cloud Volume Services for Azure/GCP). The manager leading that team was technically strong but struggling with prioritization and stakeholder management across time zones.

Task: I needed to address the delivery gap without losing a technically valuable leader, while also ensuring we met our commitments to the Azure and GCP partnerships — which were critical to NetApp's cloud strategy.

Action: I had a direct, data-driven conversation with the manager, showing the pattern of missed milestones and its downstream impact. Rather than a punitive approach, I diagnosed the root cause: the manager was trying to solve every technical problem personally rather than empowering the team. I paired them with a more experienced peer for mentorship on delegation, restructured their scope to reduce context-switching, and established weekly checkpoint reviews with clear escalation criteria. I also set a 90-day improvement timeline with specific, measurable outcomes.

Result: Within 90 days, the team's delivery cadence normalized. We successfully launched the Storage Cloud Infrastructure (PaaS) deployments for Azure and GCP on schedule. The manager developed stronger delegation skills and went on to lead a larger scope. This experience reinforced my philosophy: invest in coaching first, but set clear expectations and timelines for improvement.

6. Cross-Functional Influence Without Authority

Q: Describe a time you had to influence teams or leaders outside your direct organization to achieve a critical outcome.

Situation: At AWS, I identified an opportunity to enable Aurora DB seeding for the Redshift ZeroETL integration — a feature that would allow customers to seamlessly move data from Aurora to Redshift without traditional ETL pipelines. This required deep collaboration between the Aurora MySQL team and the Redshift organization, which had different priorities and timelines.

Task: I needed to align two separate service teams (Aurora and Redshift) on a shared technical approach, agree on API contracts, and coordinate delivery timelines — all without having authority over the Redshift team.

Action: I built the business case showing customer demand and revenue potential for the combined offering. I initiated joint architecture reviews, proposed a Parallel Export mechanism (which my team owned) as the data movement layer, and negotiated shared milestones. I escalated blockers through appropriate channels while maintaining a collaborative relationship. I also ensured my team delivered the Aurora-side components ahead of schedule to build trust and reduce risk for the partner team.

Result: We successfully delivered Aurora DB seeding for Redshift ZeroETL, creating a seamless data integration path for customers. The Parallel Export feature delivered 100x speed improvement, making the integration performant at scale. This became a flagship cross-service integration for AWS, demonstrating how Aurora could serve as a data source for the broader analytics ecosystem.

7. Strategic Vision and Roadmap

Q: How do you set technical strategy and prioritize across competing demands for a large organization?

Situation: As Senior Director at NetApp HCI, I owned a \$250M software product portfolio with competing demands: maintaining the \$2B ONTAP revenue base, investing in cloud-native services (CVS for Azure/GCP), building the HCC management plane, and modernizing the SRE operations model — all with a finite engineering budget across 180+ SDEs.

Task: I needed to create a coherent multi-year strategy that balanced sustaining revenue, cloud transformation, and operational excellence — while making hard trade-off decisions about where to invest and where to divest.

Action: I developed a three-horizon framework: Horizon 1 (sustain ONTAP revenue with predictable releases and 5-nines availability), Horizon 2 (build cloud services — CVS, Cloud ONTAP, HCC), and Horizon 3 (invest in ML/AI for predictive operations). I allocated engineering capacity roughly 50/30/20 across these horizons. I established SMART telemetry analytics to predict and prevent 80% of cluster faults — reducing reactive support costs and freeing capacity for cloud investments. I made the difficult decision to automate away manual SRE tasks rather than growing headcount, investing in tooling that would compound over time.

Result: The portfolio delivered \$2B in ONTAP maintenance revenue while simultaneously launching cloud services across Azure and GCP. SMART telemetry prevented 80% of cluster faults, reducing support costs. The CI/CD left-shift reduced field regressions by 90%. The ML automation service saved \$3M. This balanced approach allowed NetApp to protect its core business while building its cloud future — and the framework I established continued to guide investment decisions after my departure.

8. Developing Talent and Succession Planning

Q: How do you think about developing your team members' careers and building a leadership pipeline?

Situation: At AWS Aurora, I managed 3 SDMs leading 50+ SDEs. With the rapid growth of Aurora MySQL (multi-billion dollar service), I needed to ensure I was developing the next generation of leaders who could take on expanded scope — particularly as we launched new initiatives like Aurora Limitless and ZeroETL integration.

Task: I needed to create a systematic approach to talent development that would produce leaders ready for promotion while maintaining delivery velocity on critical features.

Action: I implemented a philosophy that someone should be performing at least 25% of the next-level role before promotion consideration. I created stretch assignments aligned to each SDM's growth areas — one SDM led the cross-org ZeroETL collaboration (developing influence skills), another owned the Limitless Beta launch (developing ambiguity tolerance), and the third drove the AI/ML Minions initiative (developing technical strategy skills). I held monthly career conversations separate from project reviews, provided direct feedback on leadership behaviors, and ensured each SDM had visibility with senior leadership through OP reviews and business presentations.

Result: All three SDMs developed measurably stronger leadership capabilities. The team successfully delivered Aurora Serverless V2 (~\$1B impact), Zero Downtime upgrades for millions of databases, and moved Limitless to Beta — demonstrating that investing in people development and delivery excellence are not mutually exclusive. The leadership pipeline I built ensured continuity and scalability for the Aurora MySQL organization.

9. Managing Through Crisis

Q: Tell me about a time you led your team through a critical incident or high-pressure situation.

Situation: At AWS Aurora MySQL, we encountered a situation where Zero Downtime upgrades needed to be rolled out across millions of Aurora database servers with less than 10 seconds of downtime per instance. Any failure in this mechanism could impact customer workloads at massive scale — this was a multi-billion dollar service with enterprise customers depending on continuous availability.

Task: I needed to orchestrate a safe, phased rollout across millions of databases while ensuring my team had clear roles, escalation paths, and rollback procedures — all under intense executive scrutiny.

Action: I mobilized all SDMs and their teams into a war-room structure, with each SDM owning the component areas they were responsible for. I established an incident command control model with clear decision authority, implemented canary-based progressive rollouts with automated rollback triggers, and set up real-time dashboards for monitoring. I personally led the executive communication cadence, providing updates at defined intervals. I also ensured the team maintained sustainable work hours by rotating on-call responsibilities across time zones.

Result: We successfully delivered Zero Downtime upgrades across millions of Aurora DB servers with less than 10 seconds of impact per instance. Zero customer-impacting incidents during the rollout. The incident command model I established became the standard operating procedure for all future large-scale Aurora deployments. The team's confidence in handling high-stakes launches increased significantly.

10. Operational Excellence and Metrics

Q: How do you measure the health of your organization and drive continuous improvement?

Situation: When I took ownership of the Aurora MySQL service at AWS, I inherited a team that was delivering features but lacked systematic operational health metrics. There was no unified view of team productivity, service reliability trends, or operational toil — making it difficult to identify where to invest for maximum leverage.

Task: I needed to establish a comprehensive metrics framework that would give me visibility into both service health and team health, enabling data-driven decisions about where to invest engineering capacity.

Action: I defined KPIs across four dimensions: Service Health (availability, latency, error rates), Delivery Velocity (feature throughput, cycle time), Operational Burden (toil hours, pages per engineer, COE frequency), and Team Health (engagement, attrition risk, promotion pipeline). I implemented automated dashboards and established weekly operational reviews where SDMs presented their metrics with root-cause analysis for any degradation. I used the operational burden metrics specifically to justify the AI/ML Minions investment — showing leadership exactly how much capacity was being consumed by toil.

Result: The metrics framework revealed that 56% of engineering capacity was consumed by operational tasks — a data point that secured investment approval for the AI/ML automation initiative. After implementation, we reclaimed that capacity for feature development. Canary optimizations (identified through the metrics) saved \$6M annually. The framework also improved team health by making workload distribution visible and actionable, reducing burnout risk across the 50+ person organization.

11. Diversity, Equity, and Inclusion

Q: How have you built diverse and inclusive teams, particularly across global organizations?

Situation: As Senior Director at NetApp HCI, I managed 180+ SDEs distributed across USA, EMEA, and APAC regions. Each region had different cultural norms around communication, decision-making, and career expectations. I noticed that APAC team members were less likely to speak up in global meetings and were underrepresented in leadership discussions.

Task: I needed to create an inclusive environment where all team members — regardless of geography or cultural background — had equal opportunity to contribute, grow, and be recognized.

Action: I restructured meeting cadences to accommodate time zones equitably (rotating meeting times rather than always favoring US hours). I established regional leadership forums where APAC and EMEA leaders could drive decisions locally rather than always deferring to US-based leadership. I ensured promotion and recognition processes explicitly considered contributions from all regions, and I personally mentored high-potential engineers from underrepresented regions. I also created cross-regional project assignments to build relationships and break down silos.

Result: APAC team members' participation in architecture reviews and leadership discussions increased measurably. We promoted leaders from all three regions into expanded roles. The cross-regional collaboration model enabled us to achieve 5-nines availability for cloud operations by leveraging follow-the-sun coverage. Team engagement scores improved across all regions, and the inclusive culture became a retention advantage in competitive talent markets.

12. Innovation and Technical Judgment

Q: Tell me about a time you made a bold technical bet that paid off for your organization.

Situation: At NetApp, I observed that our support organization was spending millions annually on reactive incident management — engineers manually triaging customer issues, often misrouting them, leading to poor customer satisfaction. Traditional rule-based automation had plateaued in effectiveness.

Task: I needed to find a fundamentally different approach to incident management that could scale with our growing customer base (8K+ HCI clusters) without linearly growing headcount.

Action: I championed the adoption of Machine Learning for incident analytics — a bold bet at the time (pre-2018) when ML in enterprise operations was not mainstream. I secured budget for an ML automation service, hired data science talent, and built ML models for incident auto-routing and performance signature detection. I personally shielded the team from short-term pressure to show ROI, arguing for a 12-month investment horizon. I also established the Active IQ Telemetry SaaS platform on AWS to provide the data foundation for ML models.

Result: The ML automation service saved \$3M directly. ML-based incident analytics saved an additional \$2M annually while improving customer satisfaction (CSAT) by 30%. Performance signatures reduced operational costs by 30%. The SMART telemetry system prevented 80% of cluster

faults proactively. This investment transformed NetApp's support model from reactive to predictive — and the Active IQ platform became a strategic asset for the company's cloud transformation.
